

ASSEMBLYAI VOICE AGENT HACKATHON • 2026

GrowthVoice OS

Every after-hours visitor gets an answer.

A voice agent for your website that qualifies visitors out loud and writes leads into your CRM mid-call — with US disclosure and consent built in.

growthvoice-os.vercel.app



Buyers show up after hours. The website answers the next morning.



9:40 pm

A buyer on your pricing page has one question.



9:42 pm

They fill in the contact form and leave.



9:15 am

Your reply arrives the next business day.



By then

They have booked a call with someone who answered.

What that buyer is offered today

Contact form

a reply tomorrow

Text chatbot

easy to ignore

Human SDR, nights and weekends

expensive to staff

One voice agent that answers, acts and stays compliant



Answers in voice

A visitor clicks and speaks. No phone number, no download — it runs in the browser.



Acts mid-call


It calls tools while it talks, so the lead is created and qualified in your CRM before the call ends.



Compliant by design

It says it is an AI up front and asks consent where state law requires it — then keeps the record.

How a call works



Audio goes from the browser straight to AssemblyAI. The API key never leaves the server.



Each tool.call runs over HTTP and returns as tool.result — so the agent keeps talking with real data.

Built on the AssemblyAI Voice Agent API

7 / 7

tools registered in session.update

Verified against the live API: all seven
echoed back in session.updated.



Ephemeral tokens

Minted server-side; the key never reaches the browser.



Direct streaming

24 kHz PCM16 straight to AssemblyAI over WebSocket.



session.update

Prompt, disclosure greeting, voice and turn detection.



tool.call → tool.result

Seven tools, executed and answered mid-conversation.



Barge-in

interrupt_response on; playback stops when interrupted.

Compliant before it listens

California AB 2905

Disclose an AI voice before any substantive exchange.

Texas SB 140

Disclose within the first 30 seconds.

All-party-consent states

Live transcription can count as interception, so ask first.

Explicit opt-in required in 13 states

CA

CT

DE

FL

IL

MD

MA

MI

MT

NV

NH

PA

WA



Unknown location: the strictest policy applies.



Consent is recorded before any audio is captured.

Priced from real cost

Starter

\$149 /mo

\$119/mo billed yearly

- ✓ 500 voice minutes / month
- ✓ \$0.30 / min after that
- ✓ 1 agent persona

RECOMMENDED

Pro

\$449 /mo

\$359/mo billed yearly

- ✓ 1,500 voice minutes / month
- ✓ \$0.25 / min after that
- ✓ Content pipeline and encrypted credentials

Enterprise

\$1,497 /mo

\$1,197/mo billed yearly

- ✓ 5,000 voice minutes / month
- ✓ \$0.20 / min after that
- ✓ Separate brands, unlimited personas



Stripe Checkout. A plan activates only when Stripe's signed webhook confirms payment.

Unit economics that hold at full usage

\$0.075 / min

AssemblyAI Voice Agent API, all-in


72%

gross margin even if every included minute is used

7.8 months

one extra closed deal pays for Pro at a \$3,500 contract

Pro gross margin by share of included minutes used



After Stripe's 2.9% + \$0.30. Overage never sells below cost, and a test in the suite enforces these floors.

What's real today — and what's next

Live now

- ✓ Voice calls on the AssemblyAI Voice Agent API
- ✓ Tool calls that write leads to Postgres
- ✓ Embeddable widget with real voice sessions
- ✓ Disclosure and consent records
- ✓ Stripe Checkout (test mode)
- ✓ Encrypted third-party credentials

Next 90 days

- 🕒 Multi-tenant login and accounts
- 🕒 Usage metering from live calls
- 🕒 Calendar integration
- 🕒 Per-platform credential checks
- 🕒 Measured latency

Try it.

growthvoice-os.vercel.app

Start a call and talk to it. Then open /widget-preview to see the one-line embed on a third-party page.

github.com/dev4-gpt/VoiceAI



GrowthVoice OS — every after-hours visitor gets an answer.